

capstone

December 27, 2025

```
[ ]: #https://github.com/louangelineobjero-personal/
→AIM_CapstoneProject_Netflix_Clustering_UnsupervisedLearning.git
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import joblib
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.feature_selection import VarianceThreshold
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans,DBSCAN,AgglomerativeClustering
from sklearn.metrics import silhouette_score

print("Project Domain : \nClustering Option (Unsupervised): Group customers, behaviours, or products based on similarity (K-Means, DBSCAN, Hierarchical)")
print("\nStep 1: Problem Understanding & Framing")
print("Problem Statement:\n\tGroup users into distinct segments based on their viewing behavior, patterns and demographics using unsupervised clustering techniques.")
print("Task Type:\n\tClustering (Unsupervised Learning): Group customers, behaviours, or products based on similarity (K-Means, DBSCAN, Hierarchical)")
print("Target Metric:\n\tSilhouette Score which measures how well each data point fits within its cluster. Higher score indicates better cluster separation.\n\tElbow Method for choosing number of clusters.")
print("Business KPIs (Impact Metrics): \n\tCustomer Engagement - increase monthly active sessions per customer (%) \n\tMarketing cost reduction - targeted and segment based campaigns (%)")

#Step 2: Data Collection & Understanding
print("\nStep 2: Data Collection & Understanding \nManually downloaded Netflix 2025 User Behavior Dataset (210K+ Records) in https://www.kaggle.com/datasets/sayeeduddin/netflix-2025user-behavior-dataset-210k-records")
print("FILES: users.csv,watch_history.csv,movies.csv\n")
#Load Data
```

```

users = pd.read_csv("/Users/louangelineobjero/Desktop/AIMS FILES/CAPSTONE/
↪archive/users.csv")
watch_history = pd.read_csv("/Users/louangelineobjero/Desktop/AIMS FILES/
↪CAPSTONE/archive/watch_history.csv")
movies = pd.read_csv("/Users/louangelineobjero/Desktop/AIMS FILES/CAPSTONE/
↪archive/movies.csv")

#Combine genres into one column
movies['genre_combined'] = movies[['genre_primary','genre_secondary']].
↪apply(lambda x: ','.join(x.dropna().astype(str)), axis=1)

#Merge watch history with movie info
watch_merged = watch_history.
↪merge(movies[['movie_id','genre_combined']],how='left',left_on='movie_id',right_on='movie_id')

#Convert timestamps to datetime
watch_merged['watch_date'] = pd.to_datetime(watch_merged['watch_date'])

#Aggregate features per user
user_features = watch_merged.groupby('user_id').agg(total_views = ('movie_id', 'sum'),
↪'count'),avg_watch_duration = ('watch_duration_minutes', 'mean'),active_days = ('watch_date', lambda x: x.nunique())).reset_index()

# Split combined genres into lists
watch_merged['genre_list'] = watch_merged['genre_combined'].str.split(',')
watch_exploded = watch_merged.explode('genre_list')

# Count how many times each genre is watched per user
genre_counts = watch_exploded.
↪pivot_table(index='user_id',columns='genre_list',values='movie_id',aggfunc='count',fill_value=0)
↪reset_index()
user_features = user_features.merge(genre_counts, on='user_id', how='left')

#Merge with user demographics
user_features = user_features.merge(users, on='user_id', how='left')

print(user_features.head())
print(user_features.info())
print(user_features.shape)
print(user_features.isnull().sum())

print("\nStep 3: Data Preprocessing, Applied EDA & Feature Engineering ")

#Clean Data
numeric_cols = user_features.select_dtypes(include=[np.number]).columns.tolist()

```

```

categorical_cols = user_features.select_dtypes(include=['object']).columns.
    ↪tolist()

# Numeric: fill with media
user_features[numeric_cols] = user_features[numeric_cols].
    ↪fillna(user_features[numeric_cols].median())

# Categorical: fill with 'Unknown'
user_features[categorical_cols] = user_features[categorical_cols].
    ↪fillna('Unknown')

#Remove duplicates
user_features = user_features.drop_duplicates(subset='user_id')

#Handle outliers (clip top/bottom 1%)
for col in ['total_views', 'avg_watch_duration', 'active_days']:
    lower = user_features[col].quantile(0.01)
    upper = user_features[col].quantile(0.99)
    user_features[col] = user_features[col].clip(lower, upper)

#Engineer features

#Encode categorical features
user_features_encoded = pd.get_dummies(user_features, columns=categorical_cols, ↪
    drop_first=True)

#Scale numeric features
scaler = StandardScaler()
user_features_encoded[numeric_cols] = scaler.
    ↪fit_transform(user_features_encoded[numeric_cols])

#domain-derived feature example
genre_cols = [c for c in user_features_encoded.columns if c not in numeric_cols ↪
    and 'user_id' not in c]
if genre_cols:
    user_features_encoded['top_genre_pct'] = user_features_encoded[genre_cols].
        ↪max(axis=1) / (user_features_encoded['total_views'] + 1e-6)

#Applied EDA (Exploratory Data Analysis)

for col in ['total_views', 'avg_watch_duration', 'active_days']:
    plt.figure(figsize=(6,4))
    sns.histplot(user_features[col], bins=30, kde=True)
    plt.title(f'Distribution of {col}')
    plt.show()

plt.figure(figsize=(12,8))

```

```

sns.heatmap(user_features[numERIC_COLS].corr(), annot=True, cmap='coolwarm')
plt.title('Numeric Feature Correlations')
plt.show()

#Feature selection:

# Remove low-variance features
selector = VarianceThreshold(threshold=0.01)
cols_to_drop = ['user_id'] if 'user_id' in user_features_encoded.columns else []

selector = VarianceThreshold(threshold=0.01)
user_features_selected = selector.fit_transform(user_features_encoded.
    ↪drop(columns=cols_to_drop))

selected_features = user_features_encoded.drop(columns=cols_to_drop).
    ↪columns[selector.get_support()]
print("Selected features:", selected_features)

#Dimensionality reduction:

#PCA
pca = PCA(n_components=0.95) # retain 95% variance
user_features_pca = pca.fit_transform(user_features_selected)
print("\n\nPCA reduced shape:", user_features_pca.shape)

plt.figure(figsize=(6,4))
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('Number of Components')
plt.ylabel('Cumulative Explained Variance')
plt.title('PCA Explained Variance')
plt.show()

#t-SNE for 2D visualization
tsne = TSNE(n_components=2, random_state=42)
user_features_tsne = tsne.fit_transform(user_features_selected)

plt.figure(figsize=(6,6))
plt.scatter(user_features_tsne[:,0], user_features_tsne[:,1], s=10, alpha=0.5)
plt.title("t-SNE 2D Projection of Users")
plt.show()

#Feature Importance : PCA loadings
pca_loadings = pd.DataFrame(pca.components_.T, index=selected_features)
top_features = pca_loadings.abs().sum(axis=1).sort_values(ascending=False)
print("Top contributing features to variance:")
print(top_features.head(10))

```

```

print("\n\nStep 4: Model Implementation")

X = user_features_pca # Use PCA-reduced features

# Elbow Method
inertia = []
K = range(2, 11)

for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(X)
    inertia.append(kmeans.inertia_)

plt.figure(figsize=(6,4))
plt.plot(K, inertia, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method for K-Means')
plt.show()

# Silhouette Scores
sil_scores = []
for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    labels = kmeans.fit_predict(X)
    sil_scores.append(silhouette_score(X, labels))

plt.figure(figsize=(6,4))
plt.plot(K, sil_scores, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Silhouette Score')
plt.title('Silhouette Scores for K-Means')
plt.show()

best_k = K[np.argmax(sil_scores)]
print("Best K (Silhouette):", best_k)

#Train Final K-Means Model
kmeans_final = KMeans(n_clusters=best_k, random_state=42, n_init=10)
user_features['cluster_kmeans'] = kmeans_final.fit_predict(X)

kmeans_silhouette = silhouette_score(X, user_features['cluster_kmeans'])
print("K-Means Silhouette Score:", kmeans_silhouette)

#DBSCAN Clustering

```

```

dbscan = DBSCAN(eps=1.5, min_samples=10)
dbscan_labels = dbscan.fit_predict(X)

user_features['cluster_dbscan'] = dbscan_labels

# Ignore noise points (-1)
if len(set(dbscan_labels)) > 1 and len(set(dbscan_labels)) < len(dbscan_labels):
    dbscan_silhouette = silhouette_score(X[dbscan_labels != -1], □
                                         →dbscan_labels[dbscan_labels != -1])
else:
    dbscan_silhouette = -1

print("DBSCAN Silhouette Score:", dbscan_silhouette)
print("DBSCAN clusters (incl noise):", np.unique(dbscan_labels))

#Hierarchical Clustering (Agglomerative)
hierarchical = AgglomerativeClustering(n_clusters=best_k, linkage='ward')
hier_labels = hierarchical.fit_predict(X)

user_features['cluster_hierarchical'] = hier_labels
hier_silhouette = silhouette_score(X, hier_labels)

print("Hierarchical Silhouette Score:", hier_silhouette)

#Cluster Visualization (t-SNE)
plt.figure(figsize=(6,6))
sns.scatterplot(
    x=user_features_tsne[:,0],
    y=user_features_tsne[:,1],
    hue=user_features['cluster_kmeans'],
    palette='tab10',
    s=15
)
plt.title("K-Means Clusters (t-SNE)")
plt.legend(title='Cluster')
plt.show()

#Model Comparison (Deliverable)
model_comparison = pd.DataFrame({'Model': ['K-Means', 'DBSCAN', □
                                             →'Hierarchical'], 'Silhouette Score': □
                                         →[kmeans_silhouette, dbscan_silhouette, hier_silhouette]})

print(model_comparison)

#Cluster Profiling (Business Insight)
cluster_profile = user_features.groupby('cluster_kmeans')[numeric_cols].mean()

```

```

print(cluster_profile)

#Reproducibility (Save Models & Artefacts)
joblib.dump(kmeans_final, "/Users/louangelineobjero/Documents/GitHub/
    ↳Netflix_Clustering_UnsupervisedLearning/models/kmeans_model.pkl")
joblib.dump(scaler, "/Users/louangelineobjero/Documents/GitHub/
    ↳Netflix_Clustering_UnsupervisedLearning/models/scaler.pkl")
joblib.dump(pca, "/Users/louangelineobjero/Documents/GitHub/
    ↳Netflix_Clustering_UnsupervisedLearning/models/pca.pkl")

model_comparison.to_csv("/Users/louangelineobjero/Documents/GitHub/
    ↳Netflix_Clustering_UnsupervisedLearning/models/model_comparison.csv", ↴
    ↴index=False)
user_features.to_csv("/Users/louangelineobjero/Documents/GitHub/
    ↳Netflix_Clustering_UnsupervisedLearning/models/user_clusters.csv", index=False)

print("\n Step 5: Critical Thinking → Ethical AI & Bias Auditing")
print("The problem aims to group Netflix users into distinct segments based on ↴
    ↴viewing behavior and engagement patterns using an unsupervised clustering.
    ↴K-Means is well-suited to solve the problem as it is fully unsupervised, ↴
    ↴produces clear and discrete segments for business interpretability and ↴
    ↴supports Silhouette & Elbow metrics. With the use of PCA to make ↴
    ↴distance-based clustering meaningful and retain 95% of the variance before ↴
    ↴clustering. This preprocessing step ensures that Euclidean distance, which ↴
    ↴K-Means relies on, remains meaningful and stable.")
print("\nModel Limitations & Risks\n")
print ("1.Data Imbalance - Heavy users dominate the dataset and can cause ↴
    ↴clusters to center around power users. With this, casual users may be ↴
    ↴underrepresented.")
print("2.Overfitting (Structural) - High-dimensional genre features and this is ↴
    ↴mitigated via VarianceThreshold,PCA (95% variance) and Silhouette-based model ↴
    ↴selection.")
print("3.Clusters indirectly encode sensitive attributes\n")

print("\nBias & Fairness Audit: Gender vs Cluster Assignment")
gender_cluster_dist = pd.crosstab(
    user_features['gender'],
    user_features['cluster_kmeans'],
    normalize='index'
)
print(gender_cluster_dist)

print("\n")
cluster_probs = pd.crosstab(user_features['gender'], ↴
    ↴user_features['cluster_kmeans'], normalize='index')

```

```

# Check ratio between groups
demographic_parity = cluster_probs.max(axis=0) / cluster_probs.min(axis=0)
print(demographic_parity)

print("\nPropose Mitigations\n")
print("\tReweighting - Assign slightly higher weights to minority gender groups\u2192when applying PCA or K-Means, so they are proportionally represented in\u2192clusters.")
print("\tThresholds - Ensure no group exceeds or falls below a certain\u2192percentage in any clusters.")
print("\tAugmentation - Apply techniques such as oversampling behavioral\u2192patterns of underrepresented users.")
print("\tPost-Processing - Implement human checks before deploying segments")

```

Project Domain :

Clustering Option (Unsupervised): Group customers, behaviours, or products based on similarity (K-Means, DBSCAN, Hierarchical)

Step 1: Problem Understanding & Framing

Problem Statement:

Group users into distinct segments based on their viewing behavior, patterns and demographics using unsupervised clustering techniques.

Task Type:

Clustering (Unsupervised Learning): Group customers, behaviours, or products based on similarity (K-Means, DBSCAN, Hierarchical)

Target Metric:

Silhouette Score which measures how well each data point fits within its cluster. Higher score indicates better cluster separation.

Elbow Method for choosing number of clusters.

Business KPIs (Impact Metrics):

Customer Engagement - increase monthly active sessions per customer (%)

Marketing cost reduction - targeted and segment based campaigns (%)

Step 2: Data Collection & Understanding

Manually downloaded Netflix 2025:User Behavior Dataset (210K+ Records) in
<https://www.kaggle.com/datasets/sayeeduddin/netflix-2025user-behavior-dataset-210k-records>

FILES: users.csv,watch_history.csv,movies.csv

	user_id	total_views	avg_watch_duration	active_days	Action	\
0	user_00001	14	71.466667	12	1	
1	user_00002	16	48.886667	14	1	
2	user_00003	9	73.433333	8	0	
3	user_00004	17	91.300000	14	0	
4	user_00005	10	80.133333	9	1	

Adventure Animation Biography Comedy Crime ... country \

```

0      4      0      2      0      1    ...    USA
1      0      0      1      0      0    ...    USA
2      0      4      1      1      0    ...    USA
3      2      1      0      0      0    ...    USA
4      3      0      1      0      0    ...    USA

      state_province          city subscription_plan \
0  Massachusetts  North Jefferyhaven           Basic
1      Texas        North Noahstad       Premium+
2  Michigan        Traciebury        Standard
3      Ohio        South Noah        Standard
4  Arizona        West Donald        Standard

      subscription_start_date  is_active monthly_spend primary_device \
0            2024-04-08     True      36.06      Laptop
1            2024-05-24     True      14.59    Desktop
2            2023-09-22    False      11.71    Desktop
3            2024-08-21     True      28.56      Laptop
4            2024-10-28     True       9.54    Desktop

      household_size          created_at
0            1.0  2023-04-01 14:40:50.540242
1            2.0  2024-10-10 15:39:11.030515
2            3.0  2024-06-29 14:27:49.560875
3            2.0  2023-04-11 01:01:59.614841
4            6.0  2025-04-12 19:59:30.137806

[5 rows x 39 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10300 entries, 0 to 10299
Data columns (total 39 columns):
 #   Column          Non-Null Count Dtype
 ---  -- 
 0   user_id         10300 non-null  object
 1   total_views     10300 non-null  int64
 2   avg_watch_duration  10300 non-null  float64
 3   active_days     10300 non-null  int64
 4   Action           10300 non-null  int64
 5   Adventure        10300 non-null  int64
 6   Animation         10300 non-null  int64
 7   Biography         10300 non-null  int64
 8   Comedy            10300 non-null  int64
 9   Crime             10300 non-null  int64
 10  Documentary       10300 non-null  int64
 11  Drama             10300 non-null  int64
 12  Family            10300 non-null  int64
 13  Fantasy           10300 non-null  int64
 14  History           10300 non-null  int64

```

```

15 Horror          10300 non-null int64
16 Music           10300 non-null int64
17 Mystery         10300 non-null int64
18 Romance          10300 non-null int64
19 Sci-Fi          10300 non-null int64
20 Sport            10300 non-null int64
21 Thriller         10300 non-null int64
22 War              10300 non-null int64
23 Western           10300 non-null int64
24 email             10300 non-null object
25 first_name        10300 non-null object
26 last_name          10300 non-null object
27 age                9071 non-null float64
28 gender             9476 non-null object
29 country            10300 non-null object
30 state_province      10300 non-null object
31 city               10300 non-null object
32 subscription_plan    10300 non-null object
33 subscription_start_date 10300 non-null object
34 is_active           10300 non-null bool
35 monthly_spend        9283 non-null float64
36 primary_device        10300 non-null object
37 household_size         8755 non-null float64
38 created_at            10300 non-null object
dtypes: bool(1), float64(4), int64(22), object(12)
memory usage: 3.0+ MB
None
(10300, 39)
user_id                  0
total_views                0
avg_watch_duration        0
active_days                 0
Action                      0
Adventure                    0
Animation                     0
Biography                     0
Comedy                       0
Crime                         0
Documentary                   0
Drama                          0
Family                        0
Fantasy                       0
History                       0
Horror                         0
Music                          0
Mystery                       0
Romance                       0
Sci-Fi                        0

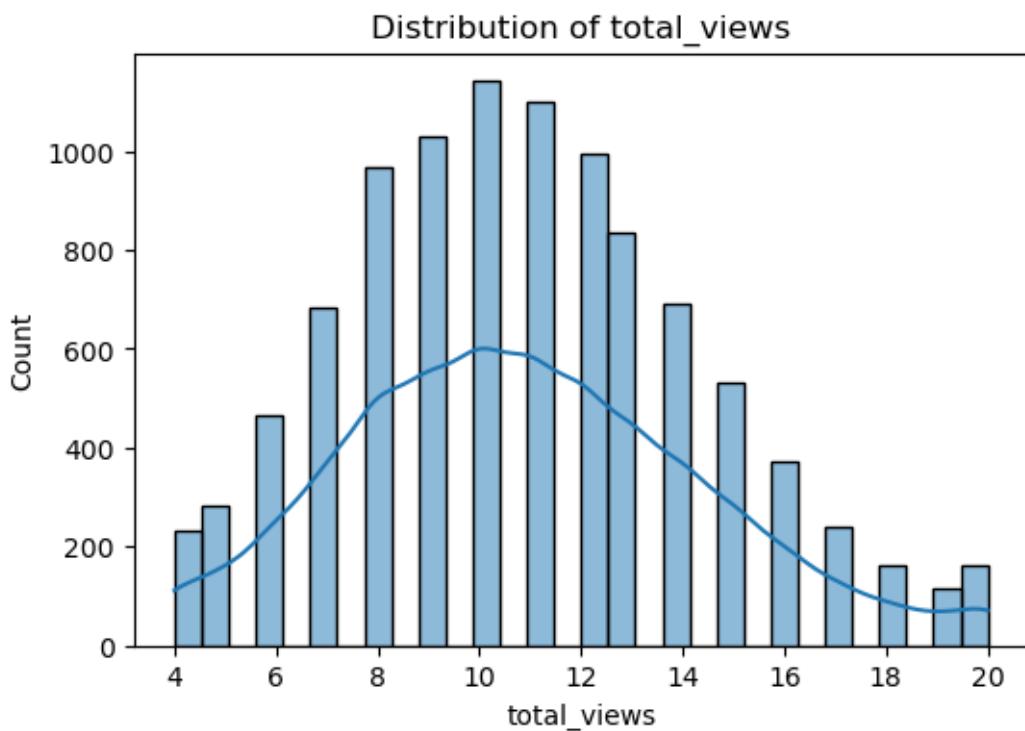
```

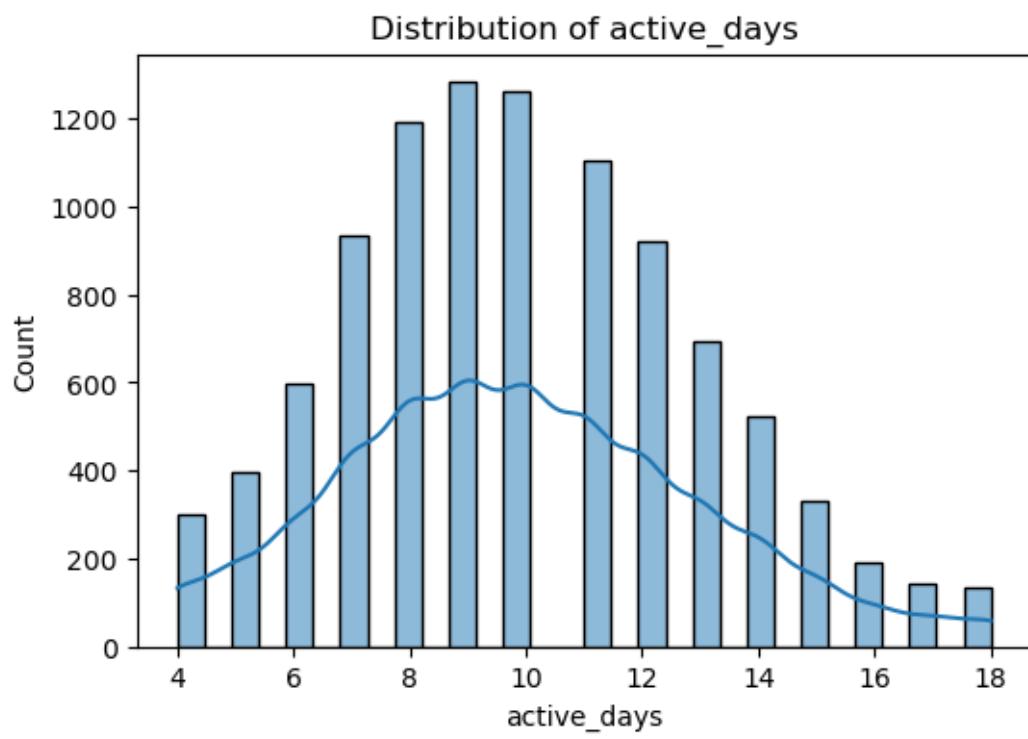
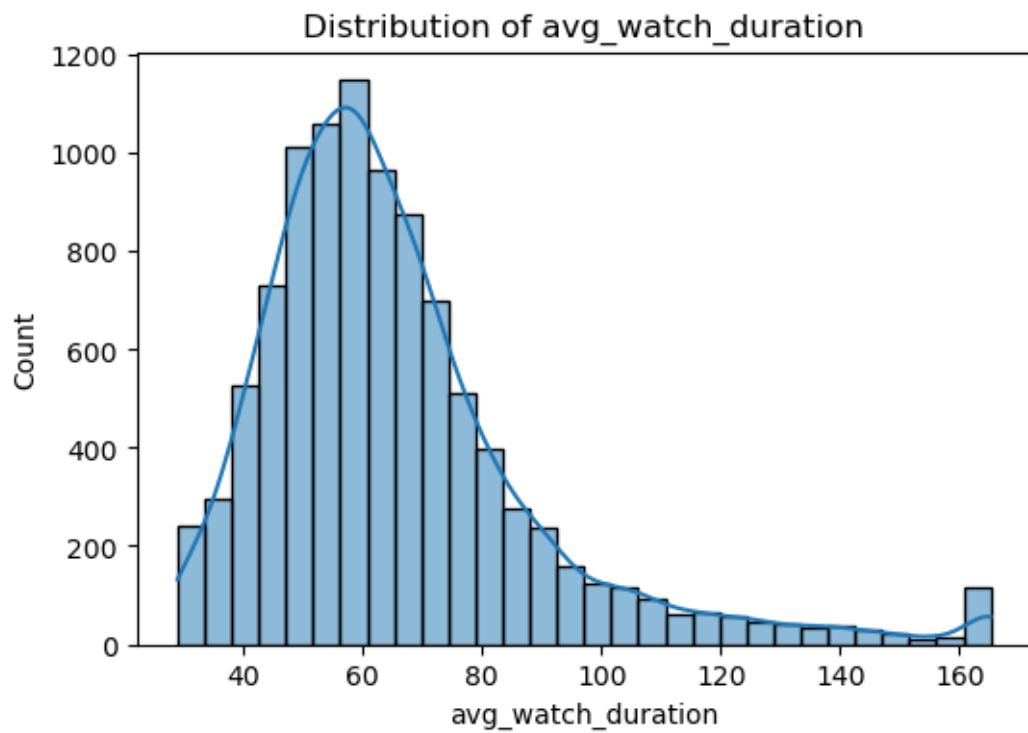
```

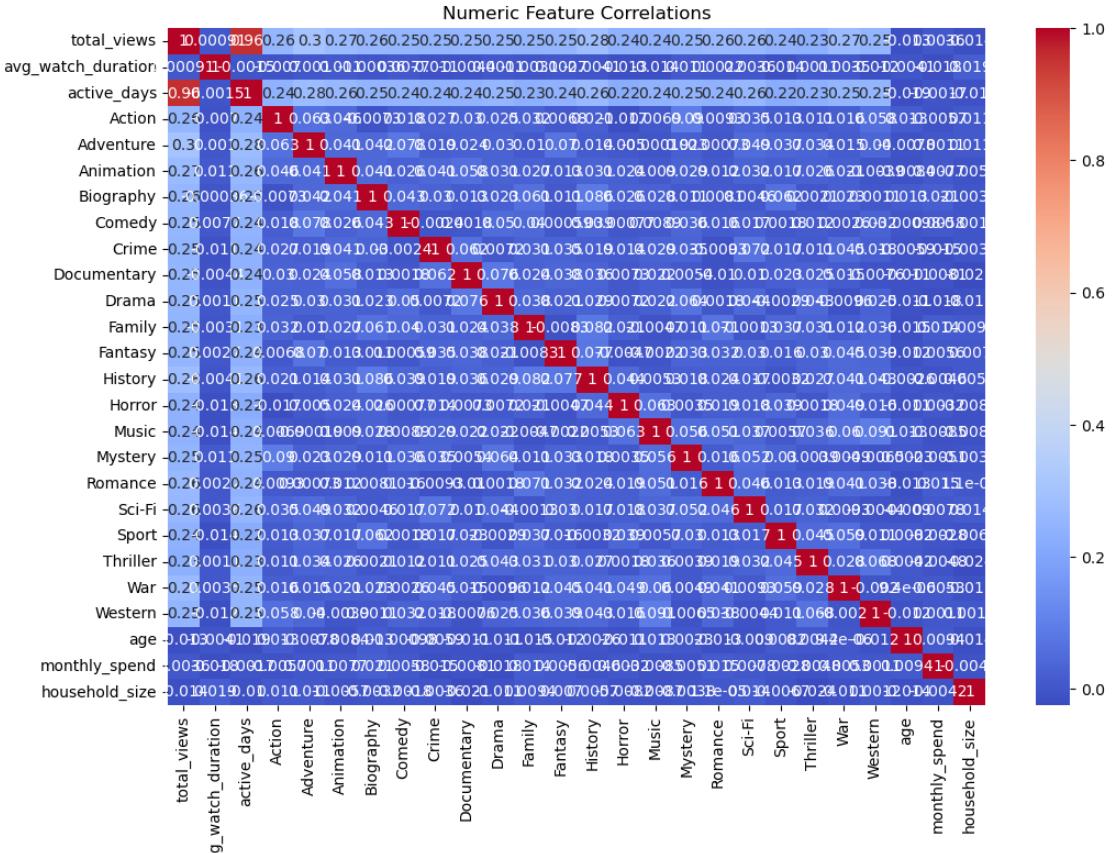
Sport          0
Thriller      0
War           0
Western        0
email          0
first_name     0
last_name      0
age            1229
gender         824
country        0
state_province 0
city           0
subscription_plan 0
subscription_start_date 0
is_active       0
monthly_spend   1017
primary_device   0
household_size    1545
created_at       0
dtype: int64

```

Step 3: Data Preprocessing, Applied EDA & Feature Engineering



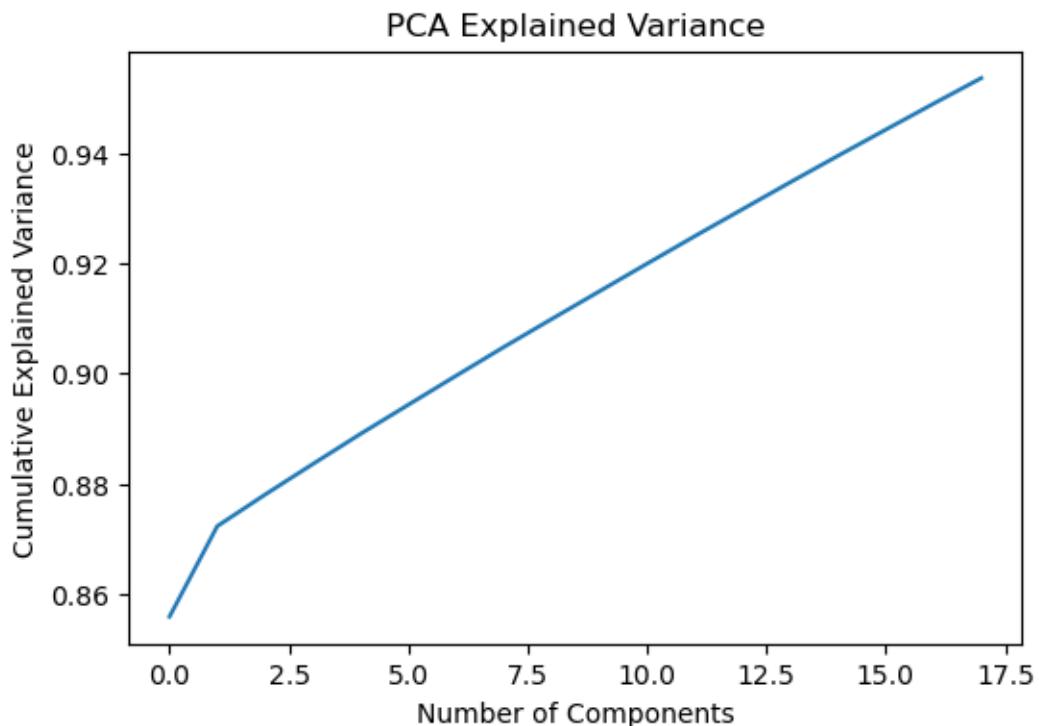




```
Selected features: Index(['total_views', 'avg_watch_duration', 'active_days', 'Action', 'Adventure', 'Animation', 'Biography', 'Comedy', 'Crime', 'Documentary', 'Drama', 'Family', 'Fantasy', 'History', 'Horror', 'Music', 'Mystery', 'Romance', 'Sci-Fi', 'Sport', 'Thriller', 'War', 'Western', 'age', 'monthly_spend', 'household_size', 'first_name_Christopher', 'first_name_David', 'first_name_James', 'first_name_Jennifer', 'first_name_Jessica', 'first_name_John', 'first_name_Matthew', 'first_name_Michael', 'first_name_Robert', 'first_name_William', 'last_name_Brown', 'last_name_Johnson', 'last_name_Jones', 'last_name_Smith', 'last_name_Williams', 'gender_Male', 'gender_Other', 'gender_Prefer not to say', 'gender_Unknown', 'country_USA', 'state_province_Arizona', 'state_province_British Columbia', 'state_province_California', 'state_province_Florida', 'state_province_Georgia', 'state_province_Illinois', 'state_province_Indiana', 'state_province_Manitoba', 'state_province_Maryland', 'state_province_Massachusetts', 'state_province_Michigan', 'state_province_Missouri', 'state_province_New Brunswick', 'state_province_New Jersey', 'state_province_New York'], dtype='object')
```

```
'state_province_Newfoundland and Labrador',
'state_province_North Carolina', 'state_province_Nova Scotia',
'state_province_Ohio', 'state_province_Ontario',
'state_province_Pennsylvania', 'state_province_Prince Edward Island',
'state_province_Quebec', 'state_province_Saskatchewan',
'state_province_Tennessee', 'state_province_Texas',
'state_province_Virginia', 'state_province_Washington',
'state_province_Wisconsin', 'subscription_plan_Premium',
'subscription_plan_Premium+', 'subscription_plan_Standard',
'primary_device_Gaming Console', 'primary_device_Laptop',
'primary_device_Mobile', 'primary_device_Smart TV',
'primary_device_Tablet', 'top_genre_pct'],
dtype='object')
```

PCA reduced shape: (10000, 18)



[]: